

Continuous-Time Radar-Inertial Odometry with B-Spline Sliding-Window Optimization

Timo Weiss and Lukas Pries

Abstract—We present a radar-inertial odometry (RIO) system for 6-DoF state estimation on a quadrotor using a single-chip 60 GHz mmWave radar (TI IWR6843, 10–60 points per frame) and a 1 kHz IMU, evaluated through aggressive racing and attitude-tumbling flight up to 10 rad s^{-1} : a regime existing RIO evaluations do not cover, where the radar model itself degrades. The trajectory is a continuous-time quintic position B-spline and a cumulative SO(3) orientation B-spline, optimized by a fixed-lag smoother with Schur complement marginalization at a 0.3 s stride over a 3 s window; a sensor-only pipeline initializes it without external *pose* reference. Doppler measures velocity directly, so we evaluate RIO as a *velocity and orientation* source, reporting position as drift per distance. A Markov-blanket factor-selection rule makes the marginalization a true marginal, and a single dynamics-adaptive measurement-weighting law covers hover to 10 rad s^{-1} , validated on held-out racing flights. At the live leading edge (the newest window estimate, before later windows refine it), velocity and orientation RMSE stay below 0.5 m s^{-1} and 3° across racing (position drift around 1 %) and degrade gracefully to 6.4° orientation (2.9 % drift) under tumbling. The output orientation covariance is approximately calibrated (NEES) on the racing flights. The orientation figures assume an *idealized* external heading reference; a perturbed one raises racing orientation error to $\sim 2\text{--}3.5^\circ$. We also document negative results: a diagnostic for marginalization inconsistency, and that the tumbling elevation bias is a flight-regime proxy, not an antenna constant.

Index Terms—radar-inertial odometry, continuous-time estimation, B-spline, sliding window, marginalization, mmWave radar, Doppler velocity

I. INTRODUCTION

Visual odometry and light detection and ranging (LiDAR)-based simultaneous localization and mapping (SLAM) are well-established, but both degrade in dusty, foggy, featureless, or textureless environments. Millimeter-wave (mmWave) frequency-modulated continuous-wave (FMCW) radar is inherently robust to these conditions: it penetrates smoke and dust, operates in the dark, and measures Doppler velocity directly without feature matching. These properties make it attractive for aggressive flight in unstructured environments.

Combining a mmWave radar with an inertial measurement unit (IMU) (radar-inertial odometry (RIO)) for ego-motion estimation has been explored with discrete-time extended Kalman filter (EKF) and graph-optimization formulations [1], [2]. However, the asynchronous nature of radar-IMU data (radar at $\sim 11 \text{ Hz}$, IMU at $\sim 1 \text{ kHz}$) motivates a continuous-time representation where the trajectory is a smooth function

evaluated at any sensor timestamp without interpolation approximation.

Research question. This work asks, end to end: *can a single-chip mmWave radar, fused with an IMU, provide usable 6-DoF state estimation, and if so, what does it take?* The starting point is what the sensor can and cannot do: per-point Doppler yields direct ego-velocity, but at 10–60 points per frame there are no bearing-stable features for position fixing, and yaw is a gauge freedom of Doppler+IMU. The system is therefore a *velocity and orientation* source whose position estimate drifts with distance, by construction rather than by defect. What it takes to turn that capability into accurate estimates, and where it ultimately fails, is the subject of the rest of this paper.

We contribute:

- *Single-chip radar under aggressive and tumbling flight*, characterizing how the radar measurement model breaks in this regime: rate-dependent intra-frame smearing, and the finding that the apparent elevation bias under tumbling is a flight-regime error proxy, not a physical antenna constant. Backflips at 10 rad s^{-1} on a single-chip sensor are, to our knowledge, unexplored.
- A *consistency-analyzed* fixed-lag smoother: a factor-set selection rule exploiting B-spline local support to make the Schur marginal exactly closed, distinct from and complementary to first-estimates Jacobian (FEJ), eliminating prior-weight tuning, with a real-time operating point per regime at minimal accuracy cost (0.3–0.5 s stride on a laptop).
- One *weighting law* from hover to 10 rad s^{-1} (with per-regime spline grids and two tumbling rescue terms), the velocity–orientation trade a measured Pareto curve, validated on *held-out racing flights* (tumbling awaits the repeated flips of Section VIII-A).
- Honest negative results: a diagnostic symptom pattern for marginalization inconsistency; a representation-level study that *disproves* the adaptive-knot-density hypothesis for this regime; ground-truth limits during aggressive maneuvers; and a normalized estimation error squared (NEES) check confirming the reported orientation covariance is consistent on racing flights.

Data and code. The complete system is released for independent reproduction. The datasets and our Ceres solver with Python bindings are available at <https://github.com/Dr1zz3l/spline-río>.

II. RELATED WORK

A recent survey [3] catalogues the rapidly growing use of mmWave radar in robotics across motion estimation, localiza-

Manuscript received Month XX, 2026; revised Month XX, 2026.

The authors are with the Department of Informatics, Technical University of Munich (TUM), 80333 Munich, Germany (e-mail: timo.weiss@tum.de; lukas.pries@tum.de).

Code and datasets are available at <https://github.com/Dr1zz3l/spline-río>.

tion, and mapping; we focus on the radar-inertial odometry thread most relevant to single-chip aerial fusion.

Radar ego-velocity estimation. Doppler-based ego-velocity estimation from mmWave radar was demonstrated by Kellner et al. [4] and extended to 3D with weighted least squares by Kramer et al. [2], whose RVE system fuses radar with IMU in a discrete-time EKF. Doer and Trommer [1] present an EKF-based RIO with online extrinsic calibration. Our work uses the same per-frame weighted least squares (WLS) ego-velocity estimator but embeds it in a continuous-time batch/sliding-window optimizer.

Continuous-time trajectory estimation. Furgale et al. [5] introduced uniform B-splines for continuous-time visual-inertial calibration. The basalt visual-inertial odometry (VIO) [6] uses cumulative B-splines on Lie groups for stereo-inertial odometry; we adapt this basalt-style cumulative SO(3) B-spline to radar-inertial fusion.

Marginalization in sliding-window estimators. OKVIS [7] and VINS-Mono [8] pioneered Schur complement marginalization for visual-inertial systems, propagating information from dropped keyframes as a dense Gaussian prior. We apply the same principle to continuous-time B-spline control points and document the practical challenges of prior scaling when the marginal's orientation and position information differ by ten orders of magnitude.

Continuous-time radar-inertial odometry. Continuous-time radar-inertial odometry was introduced by Ng et al. [9], who fuse Doppler measurements from *multiple automotive* radars with an IMU as a least-squares problem over a B-spline trajectory, evaluated on passenger vehicles. Burnett et al. [10] use a Gaussian-process motion prior for spinning-radar- and lidar-inertial odometry, and Štironja et al. [11] model the IMU trajectory continuously inside an RIO estimator with online spatio-temporal calibration. Yang et al. [12] (Go-RIO) similarly fuse radar velocity with an IMU in continuous time, but attack the same radar elevation drift we target (Section VII-E) with an uncertainty-aware ground-plane model their denser 4D cloud supports, rather than outlier rejection; that anchor assumes a persistently visible floor, whereas our tumbling regime sees it only intermittently, so we close the gap with a random sample consensus (RANSAC) front-end needing no ground return (a floor-plane factor gated on the frames that do observe it is future work, Section IX). Single-chip radar-inertial estimation has itself received recent attention: Huang et al. [13] exploit radar cross-section and Doppler physics to make a sparse single-chip cloud usable in a discrete-time sliding-window optimization, evaluated on ground-robot and handheld platforms. Our setting differs from all of these in sensor and regime: a *single-chip* radar (10–60 points per frame, 12-channel single-chip array) on a quadrotor in aggressive racing and attitude-tumbling flight up to 10 rad s^{-1} , a regime in which, as Section VIII-C shows, the radar measurement model itself degrades in ways invisible on ground vehicles and in benign aerial navigation.

Single-chip RIO on aerial platforms. Single-chip radar-inertial fusion has been shown on multirotors, but only for benign flight. Michalczyk et al. [14] use a discrete-time EKF with scan-to-scan point matching and Doppler (a factor-graph

reformulation follows in [15]); Girod et al. [16] run a real-time onboard M-estimator smoother on a quadrotor in cities and forests, but cap velocity at 2 m s^{-1} and anchor the vertical channel with a barometer. None uses a continuous-time spline, addresses aggressive or tumbling flight, or, as we do, recovers the vertical channel from radar and IMU alone (Section VII-E).

Sliding-window marginalization and consistency. That naïve sliding-window marginalization is inconsistent is long established in discrete time [17], where FEJ methods fix linearization points to preserve the observability nullspace. In continuous time, CLIC [18] and LIO-MARS [19] marginalize spline knots leaving the window by adjacency to the dropped knots; this is correct when the estimated states are all local to the window, but not when a *global* shared IMU bias couples every IMU factor and tempts the adjacency rule to pull in far-edge factors, producing an over-confident prior. Our contribution, narrower than and complementary to FEJ, is a factor-*set* selection rule that exploits B-spline local support to keep the marginal exactly closed despite this shared bias, together with a diagnostic symptom pattern for the inconsistent variant; both are developed in Section V-E and Section VIII-D.

Positioning against the EKF-RIO baselines. In the discrete-time UAV line, the EKF-RIO family of Doer and Trommer [1] defines the state of practice. Absolute root-mean-square error (RMSE) is not comparable across papers: our own quantization ablation (Section VII-D) shows firmware configuration alone changes Doppler resolution by $12\times$, so rather than quote their numbers we cross-validate against this line directly: Section VII-E runs our system on the public ICINS-2021 flight datasets of [20] and re-evaluates their EKF estimators under our causal metrics on the same data, and reports the regime-dependent result of porting their filters to our platform.

III. PRELIMINARIES

This section fixes the two mathematical tools the rest of the paper builds on: uniform B-splines as a continuous-time trajectory representation, both in $\mathbb{R}^3$ and, in cumulative form, on SO(3); and the maximum-a-posteriori (MAP) estimator that fits them to the measurements as a weighted nonlinear least-squares problem. A reader already fluent in basalt-style splines and factor-graph MAP estimation can skip to Section IV.

A. Uniform B-splines in $\mathbb{R}^3$

An order- N (degree $N-1$) uniform B-spline with control points $\{\mathbf{P}_i\} \subset \mathbb{R}^3$ and constant knot spacing Δt represents a curve as a weighted sum of basis functions,

$$\mathbf{p}(t) = \sum_i \mathbf{P}_i B_{i,N}(t), \quad (1)$$

where the $B_{i,N}$ are the Cox–de Boor basis functions [5]. Two properties make this representation attractive for asynchronous sensor fusion.

Local support. At any instant only N consecutive control points have nonzero basis, so on the segment $t \in [t_i, t_{i+1})$,

with normalized time $u = (t - t_i)/\Delta t \in [0, 1]$, the curve is a fixed blending of N control points,

$$\mathbf{p}(t) = \sum_{j=0}^{N-1} B_j(u) \mathbf{P}_{i+j}, \quad (2)$$

$$[B_0(u) \cdots B_{N-1}(u)] = \mathbf{u}^\top \mathbf{M}^{(N)},$$

with $\mathbf{u} = [1, u, \dots, u^{N-1}]^\top$ and a constant $N \times N$ blending matrix $\mathbf{M}^{(N)}$ [6]. Local support is what bounds each measurement's Jacobian to N knots and is the property the marginalization rule of Section V-E (Proposition 1) exploits.

Analytic derivatives. Differentiating the basis in (2) yields the velocity $\dot{\mathbf{p}}(t)$ and acceleration $\ddot{\mathbf{p}}(t)$ in closed form at any t , as lower-order splines over the same control points, with no finite differencing or inter-sample interpolation. This is what lets the accelerometer factor (which needs $\ddot{\mathbf{p}}$) and the min-snap regularizer (which needs $\mathbf{p}^{(4)}$) be evaluated exactly at each measurement's timestamp. We use a quintic ($N=6$) position spline so that snap is well defined and continuous.

B. Cumulative SO(3) B-splines

Orientation lives on the manifold SO(3), on which a weighted sum of rotations $\sum_i \mathbf{R}_i B_i$ is not itself a rotation, so (1) does not transfer directly. The *cumulative* parameterization of Sommer et al. [21], used in the basalt VIO [6], instead writes the orientation as a base rotation followed by a product of incremental rotations, each mapped back onto the manifold by the exponential map,

$$\mathbf{R}(t) = \mathbf{R}_i \prod_{j=1}^{N-1} \text{Exp}(\tilde{B}_j(u) \boldsymbol{\Omega}_j), \quad (3)$$

$$\boldsymbol{\Omega}_j = \text{Log}(\mathbf{R}_{i+j-1}^{-1} \mathbf{R}_{i+j}) \in \mathfrak{so}(3),$$

where the increments $\boldsymbol{\Omega}_j$ are the relative rotations between consecutive *absolute* knots $\mathbf{R}_i \in \text{SO}(3)$ (the optimization variables, basalt convention), and the cumulative basis is the tail sum of the ordinary basis of (2),

$$\tilde{B}_j(u) = \sum_{l=j}^{N-1} B_l(u). \quad (4)$$

Because each $\text{Exp}(\cdot)$ returns a rotation, $\mathbf{R}(t) \in \text{SO}(3)$ holds by construction for any knots. This form keeps the intuitive absolute control mesh, preserves the N -knot local support of the $\mathbb{R}^3$ case exactly (again used by Proposition 1), and yields the body angular velocity $\boldsymbol{\omega}_b(t)$ in closed form by differentiating the product (basalt's `evaluate_lie()`), with no numerical differentiation. We use a cubic ($N=4$) orientation spline; (7) in Section V-A is this construction with $N=4$ (base knot $\mathbf{R}_{k-3}$, three increments).

C. MAP Estimation as Weighted Least Squares

The control points and biases above form the finite-dimensional state $\mathcal{X}$ (Section IV-A). Each measurement $\mathbf{z}_i$ relates to the state through a known model $\mathbf{h}_i$, $\mathbf{z}_i = \mathbf{h}_i(\mathcal{X}) + \mathbf{n}_i$, with the noise $\mathbf{n}_i$ zero-mean Gaussian of covariance $\boldsymbol{\Sigma}_i$ and independent across measurements. The likelihood

is then $p(\mathbf{z} \mid \mathcal{X}) \propto \prod_i \exp(-\frac{1}{2} \mathbf{r}_i^\top \boldsymbol{\Sigma}_i^{-1} \mathbf{r}_i)$ with residuals $\mathbf{r}_i = \mathbf{z}_i - \mathbf{h}_i(\mathcal{X})$, so its negative logarithm is a weighted sum of squared residuals: maximum-likelihood estimation minimizes $\frac{1}{2} \sum_i \mathbf{r}_i^\top \boldsymbol{\Sigma}_i^{-1} \mathbf{r}_i$, in which each residual is weighted by its inverse covariance (information) $\boldsymbol{\Sigma}_i^{-1}$. For a scalar residual this is a scalar weight $\lambda_i = 1/\sigma_i^2$: a more certain sensor is trusted more. The per-stream weights $\lambda_a, \lambda_g, \lambda_\psi$ of Section V-B, and the rate-adaptive law that sets them, are precisely these information weights.

Prior knowledge $p(\mathcal{X})$, here trajectory smoothness and, online, the marginalization prior summarizing past data, turns the estimate into a MAP one, $\arg \max_{\mathcal{X}} p(\mathbf{z} \mid \mathcal{X}) p(\mathcal{X})$; its negative log contributes the extra quadratic terms E_{reg} and E_{prior} . Radar Doppler outliers are heavy-tailed and break the Gaussian assumption, so the quadratic on the Doppler residual is replaced by a Huber kernel ρ_δ (quadratic for $|\mathbf{r}| < \delta$, linear beyond), which bounds each outlier's influence while staying Gaussian-optimal on the inliers, a robust M-estimator. Collecting all terms gives the weighted nonlinear least-squares objective instantiated for our sensors in (6). We minimize it with Levenberg-Marquardt: each iteration linearizes the residuals about the current estimate, $\mathbf{r}_i(\mathcal{X} \boxplus \boldsymbol{\delta}) \approx \mathbf{r}_i + \mathbf{J}_i \boldsymbol{\delta}$, and solves the damped normal equations $(\sum_i \mathbf{J}_i^\top \mathbf{W}_i \mathbf{J}_i + \mu \mathbf{I}) \boldsymbol{\delta} = -\sum_i \mathbf{J}_i^\top \mathbf{W}_i \mathbf{r}_i$ for the on-manifold update $\boldsymbol{\delta}$, with $\mathbf{W}_i$ the (Huber-reweighted) information and μ the damping. Here $\boxplus$ denotes the generalized addition that applies a small update to each part of the state in its own geometry: ordinary addition for the vector components, and $\mathbf{R} \mapsto \mathbf{R} \text{Exp}(\boldsymbol{\delta})$ for the rotations, so the updated estimate stays a valid rotation. We supply analytic Jacobians $\mathbf{J}_i$. The continuous-time representation closes the loop here: every $\mathbf{J}_i$ needs the spline value and its derivatives at the measurement's exact timestamp, all available in closed form from Section III-A and Section III-B.

IV. SYSTEM OVERVIEW

A. Problem Formulation

Estimation problem. We estimate the continuous-time 6-DoF body trajectory (world-frame position $\mathbf{p}_w(t) \in \mathbb{R}^3$ and orientation $\mathbf{R}(t) \in \text{SO}(3)$) together with the constant IMU biases $\mathbf{b}_a, \mathbf{b}_g$, from asynchronous measurements over the estimation interval (the full trajectory in batch mode, a sliding window online): radar Doppler velocities (~ 11 Hz), accelerometer and gyroscope samples (1 kHz), and a low-rate heading prior. We adopt a *continuous-time* representation because the sensors are asynchronous and run at disparate rates: a trajectory defined for every t is evaluated at each measurement's exact timestamp, and its derivatives (velocity $\dot{\mathbf{p}}_w$, acceleration $\ddot{\mathbf{p}}_w$, and body angular velocity $\boldsymbol{\omega}_b$) follow in closed form, so no inter-sample interpolation or rigid-keyframe approximation is required.

Parameterization. The trajectory is a B-spline (Section V-A): the position is a quintic uniform B-spline with control points $\{\mathbf{P}_i\} \subset \mathbb{R}^3$, and the orientation a cumulative SO(3) B-spline with *absolute* rotation knots $\{\mathbf{R}_i\} \subset \text{SO}(3)$ (basalt convention [6]). This renders the otherwise infinite-

dimensional trajectory finite-dimensional; the control points and the IMU biases form the estimated state

$$\mathcal{X} = \{ \mathbf{P}_i, \mathbf{R}_i, \mathbf{b}_a, \mathbf{b}_g \}. \quad (5)$$

Only the radar pitch is observable among the extrinsics; the sliding window fixes it at the measured value (self-calibration analyzed in Section VIII-C).

Estimator. Treating the measurement noise as zero-mean and independent, the MAP estimate is the minimizer of the negative log-posterior, a weighted nonlinear least-squares problem over the measurement residuals, the spline regularizers, and the priors:

$$\begin{aligned} \hat{\mathcal{X}} = \arg \min_{\mathcal{X}} & \sum_k \rho_\delta(r_{D,k}^2) + \lambda_a \sum_m \|r_{a,m}\|^2 + \lambda_g \sum_m \|r_{g,m}\|^2 \\ & + \lambda_\psi \sum_n r_{\psi,n}^2 + E_{\text{reg}}(\mathcal{X}) + E_{\text{prior}}(\mathcal{X}), \end{aligned} \quad (6)$$

where r_D, r_a, r_g, r_ψ are the radar-Doppler, accelerometer, gyroscope, and heading residuals of Section V-B; ρ_δ is the Huber kernel ($\delta=1.0$ m/s) that bounds the influence of heavy-tailed Doppler outliers; the $\lambda_\bullet$ are the (rate-adaptive) information weights of Section V-B; E_{reg} collects the position min-snap and orientation smoothness regularizers (Section V-C); and E_{prior} holds the boundary and, online, the marginalization priors (Section V-E). We minimize (6) with Levenberg–Marquardt (Ceres) using analytic Jacobians. The fixed-lag estimator solves it over a 3 s window and, on each stride, marginalizes the states leaving the window into E_{prior} via the Schur complement (Section V-E), keeping each window’s problem size bounded.

B. Hardware Platform

The system runs on a quadrotor with a TI IWR6843AOPEVM mmWave radar (60–64 GHz FMCW), a Pixhawk flight controller IMU, and a Vicon motion capture (MoCap) system used for ground truth evaluation. The radar is mounted upside-down; the extrinsic rotation is $[180^\circ, 27.5^\circ, 0^\circ]$ (roll, pitch, yaw) and translation $[0.08, 0.02, -0.01]$ m in the body frame, with the 27.5° pitch the as-built downtilt measured by inclinometer (Section VII). The radar boresight lies along the body forward axis, so the Euler pitch is numerically the physical boresight downtilt. The radar operates at ~ 11 Hz with 0.049 m/s Doppler resolution (best-velocity firmware configuration). Raw IMU rate is ~ 993 Hz; both modalities are used at full rate.

C. Pipeline

Fig. 1 shows the pipeline. A three-stage sensor-only initialization (P1–P3, detailed in Section V-D) builds the initial trajectory the nonlinear solve needs, from radar and IMU data alone; the same raw measurements then also enter the solver directly as factors (Fig. 1, dashed). MoCap supplies only the yaw heading prior (a noise-free, idealized magnetometer surrogate; Section VIII-B bounds the real-sensor gap) and the final RMSE evaluation; P1–P3 and the position/attitude estimate use no external pose reference.




Fig. 1. System pipeline. Solid: initialization flow (P1–P3 seed the optimizer, run once). Dashed: raw sensor data fed as measurement factors. Blue: the sliding-window re-solve at each stride. Radar returns first pass a RANSAC ego-velocity prefilter (Section V-B) that rejects elevation-biased outliers before both P2 and the solver. MoCap supplies only the yaw heading prior (magnetometer surrogate) and RMSE evaluation.

V. METHODOLOGY

A. State Representation

Position. The world-frame position $\mathbf{p}_w(t) \in \mathbb{R}^3$ is a quintic (order-6) uniform B-spline with knot spacing $\Delta t_p = 5$ ms (slow racing; per-regime grids in Table II). Velocity $\dot{\mathbf{p}}_w$ and acceleration $\ddot{\mathbf{p}}_w$ are obtained as analytical derivatives of the spline.

Orientation. We use the *cumulative* SO(3) B-spline formulation [6], [21]:

$$\mathbf{R}(t) = \mathbf{R}_{k-3} \prod_{j=k-2}^k \text{Exp}(\tilde{B}_j(t) \text{Log}(\mathbf{R}_{j-1}^{-1} \mathbf{R}_j)), \quad (7)$$

where k is the active knot index, the $\mathbf{R}_i \in \text{SO}(3)$ are *absolute* rotation knots (the optimization variables; basalt convention [6]), $\text{Log}(\mathbf{R}_{j-1}^{-1} \mathbf{R}_j)$ are their relative increments, and $\tilde{B}_j(t)$ are cumulative cubic basis functions. Knot spacing is $\Delta t_o = 8$ ms (slow racing; see Table II). Body-frame angular velocity $\boldsymbol{\omega}_b(t)$ is obtained as a closed-form derivative of (7) directly from `evaluate_lie()` without numerical differentiation.

Biases and extrinsics. A constant 6-DoF IMU bias $[\mathbf{b}_a, \mathbf{b}_g]$ is estimated jointly. Of the radar extrinsics only the pitch δ_{pitch} is observable from Doppler (roll and yaw are fixed); the sliding window fixes it at the measured 27.5° , with batch self-calibration analyzed in Section VIII-C.

B. Measurement Models

Radar Doppler with lever arm. The TI IWR6843 reports positive Doppler for a receding target. With antenna offset $\mathbf{t}_{bs}$ from the body center of mass (CoM), the predicted Doppler for a point at sensor-frame bearing $\hat{\mathbf{u}}_s$ is:

$$v_D = -\hat{\mathbf{u}}_b^\top (\mathbf{R}_{bw} \mathbf{v}_w(t) + \boldsymbol{\omega}_b(t) \times \mathbf{t}_{bs}), \quad (8)$$

where $\hat{\mathbf{u}}_b = \mathbf{R}_{bs} \hat{\mathbf{u}}_s$, $\mathbf{R}_{bs}$ is the rotation from radar sensor to body frame (extrinsic $[\text{roll}=180^\circ, \text{pitch}=27.5^\circ, \text{yaw}=0^\circ]$; measured, Section VII), and $\mathbf{R}_{bw} = \mathbf{R}_{wb}^\top$. The leading minus makes the convention self-consistent: a sensor translating toward a static point ($\hat{\mathbf{u}}_b^\top \mathbf{R}_{bw} \mathbf{v}_w > 0$) yields $v_D < 0$, an approaching rather than receding target, as required. The residual $r_k = v_{D,\text{meas}} - v_{D,\text{pred}}$ uses Huber loss with $\delta = 1.0$ m/s, robust as defense-in-depth to any residual Doppler

noise on the surviving inlier returns; the 0.049 m s^{-1} Doppler resolution sets a lower bound but does not drive the choice. This δ is deliberately large relative to the $0.07\text{--}0.08 \text{ m/s}$ ego-velocity residuals we ultimately reach (Section VII-E): after the RANSAC prefilter the inliers are already clean, so δ caps only the rare survivor rather than shaping the fit.

RANSAC ego-velocity prefilter (default front-end). A single-chip radar's small 12-channel virtual aperture has coarse elevation resolution, and a minority of returns carry an elevation-correlated Doppler bias that a robust kernel only down-weights. Following reve [1] (Doer's radar ego-velocity estimator), we therefore prepend a reve-style 3D least-squares RANSAC ego-velocity estimate to the per-frame Doppler set (inlier threshold 0.15 m/s) and pass only the inlier returns to the spline solver, hard-rejecting the elevation-biased minority before the optimization. The prefilter is seeded (`numpy.default_rng(0)`), so the full pipeline reproduces bit-identically; frames with fewer than five returns bypass it unchanged. On those sparse frames (most of the backflip regime) the in-solver Huber loss is the sole radar outlier protection, so the two stages are complementary. It runs once at frame load, not per window, so it does not enter the per-window timing budget (Section VII). Section VII-E quantifies what it removes: on the public ICINS flights it closes an order-of-magnitude vertical-drift gap, and on our own flights it improves aggressive-racing live position by 8% ($0.51 \rightarrow 0.47 \text{ m live}$). It is enabled by default for every result in this paper.

Dynamics-adaptive measurement weighting. Our central practical finding is that a single *rate-adaptive trust allocation* across the three measurement streams covers hover to 10 rad s^{-1} with one configuration, each component derived from a measured failure diagnosis (Section VIII-C), not tuned ad hoc:

- *Gyroscope: trust more when rotating fast.* We scale the per-sample gyro weight as $\lambda_g^{\text{eff}} = \lambda_g (1 + (|\mathbf{z}_g|/\omega_g)^p)$ with $\lambda_g=4$, $\omega_g=4 \text{ rad/s}$, $p=4$, using the *measured* rate $\mathbf{z}_g$ (causal, no spline evaluation). The quartic exponent keeps racing untouched ($\lambda_g^{\text{eff}} \approx \lambda_g$ below 3 rad/s) and reaches $\lambda_g^{\text{eff}} \approx 160\text{--}500$ during flips, where the gyroscope is the only clean sensor. A quadratic ramp ($p=2$) stiffens the gyro already at $2\text{--}5 \text{ rad/s}$ (where radar is still informative) and degrades racing roll/yaw.
- *Radar: trust less when rotating fast.* A radar frame integrates over tens of milliseconds; at body rate ω the scene smears by $\omega T_{\text{frame}} (\approx 17^\circ$ at 10 rad/s). We inflate the per-frame radar noise as $\sigma_{\text{eff}}^2 = \sigma_0^2 (1 + (|\omega|/\omega_r)^2)$ with $\omega_r=4 \text{ rad/s}$ and $|\omega|$ from the warm-start spline at problem-build time; the down-weighting is continuous, with no hard-gate cliff.
- *Accelerometer: same form, $\omega_a=8 \text{ rad/s}$:* during flips the accelerometer carries thrust transients and vibration, its gravity information is nil mid-flip, and it otherwise distorts orientation through the $\mathbf{R}\text{--}\dot{\mathbf{p}}$ coupling; $(\omega_r, \omega_a)=(4, 8)$ is the knee of a measured velocity–orientation trade (Section VIII-C).
- *Per-point signal-to-noise ratio (SNR) weighting.* $w_i = \text{clip}((I_i/\bar{I})^\alpha, 0.25, 4)$ with frame-median intensity $\bar{I}$ and

$\alpha=1$ (median-relative, so the global radar weight is unchanged). High-SNR returns carry measurably better Doppler: on aggressive racing this alone improves live orientation from 3.31° to 2.88° at unchanged position.

Flip-regime Doppler correction. Under sustained tumbling an additional attitude-correlated Doppler systematic appears, which we absorb with a per-point term $v_{\text{corr}} = v_{\text{meas}} - b u_{s,z}$ ($b=-1.5 \text{ m/s}$, backflips only; $u_{s,z}$ is the vertical component of the sensor-frame bearing $\hat{\mathbf{u}}_s$). Section VIII-C shows why b must be interpreted as a flight-regime error proxy rather than a physical elevation bias of the array.

Accelerometer.

$$\mathbf{r}_a = \mathbf{z}_a - \mathbf{R}_{bw}(\ddot{\mathbf{p}}_w - \mathbf{g}_w) - \mathbf{b}_a, \quad (9)$$

with $\mathbf{z}_a$ the measured specific force, weighted $\lambda_a = 0.01$: this 400:1 gyro/accel ratio ($\lambda_g=4$, below) is a deliberate design statement that leans position and velocity on radar Doppler plus min-snap, since the accelerometer carries thrust transients and frame vibration that corrupt the $\mathbf{R}\text{--}\dot{\mathbf{p}}$ coupling (the dynamics-adaptive law of Section V-B down-weights it further mid-flip).

Gyroscope.

$$\mathbf{r}_g = \mathbf{z}_g - \boldsymbol{\omega}_b(t) - \mathbf{b}_g, \quad (10)$$

weighted $\lambda_g = 4.0$ (L2 loss).

Heading prior. Yaw is unobservable from Doppler alone (all yaw-equivalent trajectories predict identical radial velocities under pure rotation about gravity). We supply a heading reference from a magnetometer (standard on most platforms; here simulated by MoCap yaw). The residual $r_\psi = \psi_{\text{est}} - \psi_{\text{ref}}$ is added at $\sim 100 \text{ Hz}$, with $\lambda_\psi = 0.6$ by default and raised to 10 on the heading-critical bags (slow racing, backflips; Table III footnote). A weight sweep shows the split is a yaw-observability/position trade. On slow racing, gentle dynamics leave yaw radar-unconstrained, so the weak default ($\lambda_\psi=0.6$) lets yaw drift to 4.3° (live orientation $2.0 \rightarrow 4.2^\circ$) while a strong prior costs almost no position. On fast racing, the rich translational excitation already holds yaw near 3.8° at $\lambda_\psi=0.6$, so raising it to 10 buys only 0.5° of yaw yet degrades live position (drift $1.0 \rightarrow 1.3\%$) through the orientation–velocity coupling; we therefore keep it light. On backflips the anchor instead stabilizes the whole attitude through the flip: dropping λ_ψ to 0.6 inflates live orientation from 6.4° to 8.7° (all three axes degrade), so it stays at 10. The per-regime weight is therefore not monotone in dynamics: gentle (slow racing) and tumbling (backflips) flight need the same strong prior at opposite rates while translation-rich fast racing needs a weak one, so it cannot key off $|\omega|$. It is thus the one weight *not* folded into the runtime-adaptive law above, set once per flight type rather than auto-weighted; keying it instead off a yaw-observability proxy such as translational excitation, and adapting it to real magnetometer noise in place of the idealized MoCap heading, is left to future work.

C. Regularization

Position: minimum snap. $E_{\text{snap}} = \lambda_s \int \|\mathbf{p}^{(4)}(t)\|^2 dt$ with $\lambda_s = 2 \times 10^{-5}$.

Orientation: minimum angular acceleration. Rather than penalizing angular velocity increments (minimum ω), which opposes every banked turn and the entire backflip maneuver, we penalize the *second* finite difference on SO(3):

$$\mathbf{r}_\alpha = \text{Log}(\mathbf{R}_{i-1}^{-1}\mathbf{R}_i) - \text{Log}(\mathbf{R}_i^{-1}\mathbf{R}_{i+1}), \quad (11)$$

which is zero for constant angular rate. Swept $\lambda_\alpha \in \{0, 0.001, 0.01, 0.1, 1.0\}$; 0.1 was best across all bags (1.0 caused backflip divergence).

D. Sensor-Only Initialization (P1–P3)

P1: Orientation. A stationary segment before flight is detected automatically; the gyroscope bias is the mean reading over this segment and the accelerometer bias a gravity-aligned scalar correction, both *sensor-only* and using no external attitude reference. Initial roll/pitch are then derived from the stationary accelerometer (gravity direction). Yaw is set to zero (gauge freedom; corrected later by the heading prior). The full orientation trajectory is obtained by forward-integrating debiased gyroscope readings: $\mathbf{R}(t + \Delta t) = \mathbf{R}(t) \text{Exp}((\mathbf{z}_g - \mathbf{b}_g)\Delta t)$.

P2: Position. Per radar frame, weighted least squares estimates the 3D ego-velocity $\mathbf{v}_{\text{wls}}$ from the RANSAC inlier returns (Section V-B). Doppler alias unwrapping uses IMU-integrated world-frame velocity to select the correct alias once at frame load (23.1% of points aliased in fast racing, 2.0% in slow racing); the unwrapped Doppler is the measurement consumed by *both* this WLS estimate and the main solver’s radar factors (Eq. 8), with the solver’s own v_{max} check kept only as a residual safety net. World-frame velocity is integrated with forward Euler to produce an initial position trajectory.

P3: Boundary priors. Position, velocity, orientation, and angular velocity boundary priors are built from the sensor-only state at $t = 0$. $\lambda_{\text{bnd}} = 1000$ for position and velocity.

E. Sliding Window with Schur Complement Marginalization

The fixed-lag smoother advances a 3 s window in 0.3 s strides. After each Ceres LM solve, the *stride zone* CPs and orientation knots are marginalized:

$$S = H_{bb} - H_{ab}^\top H_{aa}^{-1} H_{ab}, \quad (12)$$

where a indexes stride-zone variables and b indexes boundary variables (last $N-1$ position CPs, last $N-1$ orientation knots, and the 6-DoF bias; dimension 30). S is factored as $S = LL^\top$ and stored as a `MargPriorFunctor` for the next window.

Re-centered prior. Before adding the prior to each window’s problem, the linearization point $\mathbf{x}_0$ is updated to the current warm-start. This makes the prior contribute only curvature (the Hessian shape), not a gradient pull toward a stale historical estimate. The residual is $\mathbf{r} = L^\top(\mathbf{x} - \mathbf{x}_0)$ in local coordinates.

Consistent factor selection (Markov-blanket rule). Which residuals enter Eq. (12) determines whether the prior is a true marginal. A naïve choice (all residuals touching marginalized or boundary variables) pulls in every IMU factor of the window through the shared bias block ($\sim 25,000$ residual rows), *conditions* on interior states (treating them as perfectly

known), and double-counts factors that are re-added in the next window. The resulting prior is overconfident by orders of magnitude and must be deflated by a hand-tuned scale of 10^{-7} – 2×10^{-4} depending on flight dynamics (Section VIII-D). The shared global bias is the trap: it is what tempts pulling in the far-edge IMU factors. We instead exploit the splines’ local support of N knots (quintic position $N=6$, cubic orientation $N=4$).

Proposition 1 (Exact closed spline marginal at a fixed linearization). *At a fixed linearization point, let $\mathcal{M}$ be the knots leaving the window and $\mathcal{B}$ the $N-1$ retained knots immediately forward of $\mathcal{M}$ plus the global 6-DoF IMU bias (the dim-30 boundary of Section V-E). Every order- N spline factor has support over N consecutive knots, and an IMU factor also touches the bias; hence any factor touching $\mathcal{M}$ has support $\subseteq \mathcal{M} \cup \mathcal{B}$ with its bias dependence in $\mathcal{B}$. Factors reaching further back were consumed in the previous marginalization step and enter the current Schur complement only through the incoming prior term, never re-added as fresh factors, so nothing is double-counted. Selecting exactly the factors touching $\mathcal{M}$ (plus that prior) therefore forms a Schur complement that conditions on no interior state and shares no factor with the next window: at this linearization it is the exact marginal.*

Including *only* residuals that touch a marginalized variable (rather than those touching marginalized or boundary variables) thus yields a closed marginal. Empirically it is also *consistent*: the prior is applied **unscaled** (scale = 1) across all flight regimes, and its Mahalanobis residual at the solution drops from 10^6 – 10^{12} to $\mathcal{O}(1)$, i.e. it agrees with the incoming data.

Relation to FEJ. This mechanism is distinct from, and complementary to, first-estimates-Jacobian approaches to marginalization consistency [17]: FEJ fixes the *linearization points* of the prior to preserve the observability nullspace, whereas our rule corrects the *factor set* entering the Schur complement so that the marginal is exactly closed at a *fixed linearization point*. We do not employ FEJ; instead the prior is re-centered at each window’s warm start (curvature retained, gradient pull discarded). We do not claim this re-centering preserves consistency across relinearizations as a theoretical guarantee; rather, in our experiments it left no measurable inconsistency symptom; the NEES analysis of Section VII-F provides the corresponding empirical evidence. Our missions are short (18–26 s), which may be too brief to expose drift a re-centered prior could accumulate over a long deployment; bounding this over extended runs is left to future work.

Live-edge warm start. Control points entering the window each stride would otherwise hold stale dead-reckoned initialization values; any CP-level seam is amplified by the min-snap term ($\propto \Delta/\Delta t_p^4$, initial cost $\sim 10^{13}$). We rigidly align the entering segment (shift + ΔR) to the last data-supported CP of the previous solution, reducing the initial cost by five orders of magnitude.

Marginalization cost. The restricted Gauss–Newton Hessian for Eq. (12) is accumulated by evaluating the affected cost functions directly (tangent-space Jacobians via the manifold

TABLE I
DATASET SUMMARY

Bag	Dur. (s)	$v_{\max}$ (m/s)	$ \omega _{\max}$ (rad/s)	Frames
Slow racing	25.8	3.65	2.62	258
Fast racing	18.0	5.66	9.14	179
Backflips*	18.1	~ 4.3 (95%)	~ 10.9 (95%)	181

*95th percentile; MoCap tracking artifacts dominate the instantaneous peaks. The ~ 10 rad/s headline is this 95th-percentile sustained rate, not an instantaneous maximum.

TABLE II
PER-REGIME SLIDING-WINDOW CONFIGURATION

Parameter	Slow	Fast	Backflips
Position knot Δt_p (ms)	5	40	10
Orientation knot Δt_o (ms)	8	16	8
LM iteration cap	12	400	400
Heading prior λ_ψ	10	0.6	10
Doppler elevation corr. b (m/s)	—	—	−1.5
Position-init anchor λ_{p_0}	—	—	0.5

Shared across all regimes: 3 s window, 0.3 s stride, unscaled Markov-blanket marginalization prior, the single dynamics-adaptive weighting law (Section V-B), RANSAC ego-velocity prefilter (0.15 m s^{−1} inlier), sensor-only initialization, and frozen extrinsic pitch 27.5°. “—” marks the shared default: uncapped-by-data iterations, and the flight-regime elevation correction / position tether off.

plus-Jacobian, Triggs correction for robust losses), avoiding a full solver re-linearization; the marginalization step costs ~ 10 ms per window.

VI. EXPERIMENTAL SETUP

All experiments use the rosbag datasets collected in a Vicon MoCap lab. Table I summarizes the three bags used for evaluation.

All three regimes run the *same* sliding-window solver, weighting law, and front-end; Table II collects the handful of per-regime knobs that differ, so that every reported number can be reproduced from one place.

IMU rate: ~ 993 Hz (full rate, no downsampling). **Evaluation:** Start-anchored rigid alignment to MoCap ground truth: the alignment rotation $\mathbf{R}_{\text{align}} = \mathbf{R}_{\text{gt},0} \mathbf{R}_{\text{est},0}^T$ and translation $\mathbf{t}_{\text{align}} = \mathbf{p}_{\text{gt},0} - \mathbf{R}_{\text{align}} \mathbf{p}_{\text{est},0}$ are fixed at frame 0 (no scale). The final 3 s is trimmed. Velocity RMSE is computed from the analytical first derivative of the position B-spline, not from finite differences of the estimated trajectory; the reference is the Vicon-fed onboard state-estimate velocity, zero-phase low-pass filtered (4th-order Butterworth, 10 Hz), not a raw finite difference of MoCap position. *Position drift* (%) is position RMSE normalised by the total GT path length, the standard odometry metric for a dead-reckoning system with no loop closure (per-segment relative pose error (RPE) in Section VII-B). Start-anchored alignment is harsher than the whole-trajectory SE(3)/Umeyama fit (the EVO/KITTI default) used by most odometry papers, because it cannot retro-fit a constant transform to cancel accumulated drift: it fixes the frame once, at the start, exactly as a deployed estimator must. We adopt it precisely for this causal, deployment-relevant reading, and apply it identically to every system compared. Where we compare against externally published figures (Section VII-E) we additionally report whole-trajectory-aligned

numbers so the baselines can be cross-checked against their own papers. For sliding window, *live* (*leading-edge*) RMSE is computed by snapshotting the estimate at $[t_{\text{end}} - \text{stride}, t_{\text{end}}]$ of each window at 50 Hz before any subsequent refinement. The *settled* RMSE evaluates the same trajectory after all windows have refined each point.

Solver implementation. The Ceres LM problem is solved with SPARSE_NORMAL_CHOLESKY (SuiteSparse), up to 400 iterations (function-tolerance stop; typically ~ 30). The analytic Jacobians combine the basalt cumulative-spline Jacobian struct for the spline terms, closed-form expressions for the IMU sensor terms, and SymForce-generated code [22] for the radar sensor term; Ceres automatic differentiation was eliminated to avoid Jet-arithmetic overhead over the large set of spline control-point parameters per factor. In the sliding-window solver the linear solve and the marginalization-prior recompute (`compute_prior`, evaluated *every* LM iteration, as distinct from the ~ 10 ms one-time prior *formation* of Section V-E) are the dominant costs; Jacobian evaluation accounts for roughly 25–30% of per-window wall time.

VII. RESULTS

A. Batch Solver and Extrinsic Pitch

We use the full-trajectory batch solve as an offline full-batch reference; offline it reaches 0.20 m / 1.1° on slow racing and 0.74 m / 2.4° on fast (the dense racing Δt_p over-fits fast racing, where the coarser sliding window does better). The radar-to-body pitch is a fixed mount angle, measured at 27–28° with an inclinometer on the radar face (reference surface 0°) and frozen at 27.5°, a few degrees below the 30° computer-aided design (CAD) design (3D-printed, hand-assembled fixture). We lock it rather than estimate it online because it is only weakly observable from Doppler: the self-calibration is init-dependent and orientation RMSE is pitch-insensitive (Section VIII-C).

B. Sliding Window: Settled vs. Live

Table III reports sliding window results (3 s window, 0.3 s stride, Schur marginalization). Fig. 2 shows trajectory overlays for all three flights; Fig. 3 shows the corresponding error timeseries.

Live-edge orientation degrades 0.4–0.6° relative to settled; this is expected because subsequent windows refine previous estimates. As a complementary per-segment view, the live-edge KITTI translational RPE converges to 1.0–2.1% beyond 20 m on racing (from 5–9% at 5 m segments, where the ~ 11 Hz radar yields few Doppler per segment) and to $\sim 5.6\%$ on backflips; full curves are released with the code. With the consistent prior, settled velocity is smooth across stride boundaries. We report all tables, plots, and the cross-validation of Section VII-E at a single *evaluation configuration* (3 s window, 0.3 s stride), chosen for best accuracy and cross-table consistency, not as a deployment claim: per-window solve time is set by window length and knot density, not the stride, and at 3 s only fast racing approaches the 0.3 s budget on a laptop CPU. Real time is nonetheless within reach: pushing the window and stride to *barely* reach it per regime costs almost

TABLE III

SLIDING WINDOW: SETTLED AND LIVE LEADING-EDGE RMSE
(CONSISTENT UNSCALED PRIOR; PER-WINDOW SOLVE TIME ON A LAPTOP CPU)

Bag	Mode	Vel (m/s)	Ori (°)	Pos (m)	Drift%	$t_{\text{win}}(\text{s})$
Slow racing	Settled	0.42	1.53	0.293	0.61	0.70
	Live edge	0.48	1.97	0.310	0.64	
Fast racing	Settled	0.25	2.34	0.382	0.85	0.35
	Live edge	0.32	2.90	0.469	1.04	
Backflips [†]	Settled	—*	5.01	1.674	3.09	0.6
	Live edge	2.35	6.35	1.545	2.85	

All bags use the unscaled consistent marginalization prior (Section V-E) and the *same* dynamics-adaptive weighting law (Section V-B). Per-regime parameters are limited to spline grids ($\Delta t_p = 5/40/10$ ms), the heading weight, and for backflips the position-init anchor ($\lambda_{\text{pos-init}} = 0.5$) and flip-regime Doppler correction $b = -1.5$ m/s (Section VIII-C). Extrinsic pitch locked at 27.5° for all bags. *Settled velocity is not reported for backflips: stride-boundary resampling injects a retrospective-evaluation artifact there, which also leaves backflips settled *position* above live (the one such inversion); the **bold** live-edge rows are the deployment-relevant figures throughout. Drift %: path lengths 48.1 m / 44.9 m / 54.1 m. Iteration-capped configurations reproduce bit-identically across repeated runs on a fixed platform/BLAS (single-thread reduction); full-convergence configurations vary by $<2\%$ (multithreaded reduction order).

nothing. Fast is real-time at the 0.3 s stride with a 2 s window, slow and backflips at a 0.5 s stride (1.5 s and 2 s windows); each drops per-window solve below its stride (0.28/0.40/0.43 s for fast/slow/backflips). Slow racing and backflips stay within run-to-run noise (backflips even improves, $6.35 \rightarrow 5.56^\circ$), and fast-racing position degrades only modestly ($1.04 \rightarrow 1.15\%$ drift). Solve times carry $\sim 20\%$ variance, so the strides are taken with margin.

Deployment asks for more than *barely* real-time. As one factor source in a larger fusion graph (Section VII-F) sharing compute with a parallel visual or lidar pipeline, a tuned, no-margin point on a dedicated CPU is a feasibility demonstration, not a deployment budget. The margin must come from the back-end, and since a banded Schur factorization is numerically unstable at our conditioning, the route is incremental smoothing, not batch re-solve: Girod et al. [16] reach comfortable onboard rates with exactly such an iSAM2 back-end (discrete-time graph), and porting our continuous-time B-spline onto it, where the shared global bias induces a Bayes-tree fill-in, is the deployment path we leave to future work. The dominant cost reductions were: the consistent prior removing the marginalization re-linearization (~ 0.6 s), coarsening Δt_p from 5 ms to 40 ms for aggressive flight (which *also* drops the LM iteration count from 28 to 11 and slightly improves orientation), and warm-start alignment removing the live-edge seam.

C. Radar Contribution

Radar is essential. Disabling the radar factors collapses the live-edge estimate to IMU dead-reckoning: position drift balloons to 5.9% on slow racing and diverges on fast racing ($\approx 99.5\%$ of distance), with velocity RMSE rising to 2.4 and 8.0 m/s as double-integrated accelerometer noise accumulates unbounded. The direct Doppler velocity constraints bound this to the Table III figures (0.6–0.9% drift, 0.32–0.48 m/s velocity), one-to-two orders of magnitude lower. Orientation improves more modestly, from $8.9^\circ/4.0^\circ$ (IMU-only) to $2.0^\circ/2.9^\circ$, the 1 kHz gyroscope already constraining it.

D. Held-Out Validation

All tuning used three bags (slow/fast racing, backflips); the final configuration was then run *unmodified* on flights never used for any decision. A second, independent fast-racing flight reaches 0.37 m/2.81° live (vs. 0.47 m/2.90° on the tuning bag), and a circular trajectory reaches 0.90 m/3.77° live. Five additional stress-tier bags from an earlier recording day use an older radar firmware (best-position firmware) with $12\times$ coarser Doppler quantization (0.604 vs. 0.049 m/s), partially unusable accelerometer axes, and per-day time offsets; the default RANSAC prefilter still keeps 46–75% of returns (no starvation), and the system degrades gracefully (1.5–7 m position), with the *old-firmware backflips* orientation reaching 8.1° settled/9.1° live (better than the 14.8° of its previous, individually tuned best) under the untuned universal configuration. The rate-adaptive weighting thus transfers across both firmware and day. Each flight condition is a single representative recording; the evaluation’s breadth comes from these held-out flights and the cross-platform validation of Section VII-E rather than repeated trials (Section VIII-A).

E. External Cross-Validation Against EKF Baselines

We cross-validate against the EKF-RIO family of Doer and Trommer [1], [20] in both directions on shared data, with their toolbox pinned at a fixed commit and every parameter deviation documented in the repository.

Our system on their data. We run our sliding-window solver, with the universal weighting configuration unmodified, on the four public ICINS-2021 indoor flight datasets [20] (TI IWR6843AOP at 10 Hz, Analog Devices ADIS16448 IMU at 409 Hz with onboard barometer, loop-closure-and-scale-corrected VINS pseudo ground truth), using their published extrinsic calibration (locked) and their $\pm 60^\circ$ elevation pre-gate. Their estimators (ekf-río, unaided yaw; ekf-yrio, Manhattan-world radar yaw aiding) are re-run from their stock configurations on the same bags, and *both* systems are evaluated with the identical causal protocol of Section VI (live estimates, start-anchored alignment, same windows and ground truth). Table IV reports the headline metrics (a) and their decomposition (b). On position our *horizontal* drift is competitive (0.14–0.30% of path, Table IV(b)); the larger total drift (0.6–1.3% vs. the baselines’ 0.3–0.6%) is almost entirely the *vertical* channel (0.63–1.28%), the single-chip elevation residual the RANSAC prefilter targets (below). Velocity RMSE 0.07–0.08 m/s matches theirs; absolute position is 0.5–1.9 m, trailing the yaw-aided ekf-yrio. The full 3-DoF orientation column favors our solver (0.51 – 1.13° vs. up to 10.1°), but this gap is almost entirely *yaw*: our heading is anchored by a pseudo-magnetometer (here a VINS pseudo-GT yaw surrogate), whereas ekf-yrio uses sensor-only Manhattan-world yaw and ekf-río is unaided. On the heading-*independent* roll and pitch the three systems are comparable (0.25 – 0.74° , Table IV(b)), so the orientation column is not an attitude win; removing our heading aiding confirms this: on flight_1 our unaided yaw drifts to 9.9° , matching unaided ekf-río (10.1°), while roll/pitch and position are essentially unchanged.




Fig. 2. Trajectory overlays for all three flights (columns): MoCap ground truth vs. the SW live edge, X-Y (top) and Y-Z (bottom) projections. Within each column both projections share a common Y range and metric scale (equal aspect, so loops keep their true shape); the scale differs between columns, so compare shapes rather than absolute sizes across columns. Start $\blacksquare$. Larger fast-racing deviations reflect higher dynamics ($\omega_{\max} = 9.1 \text{ rad/s}$); on backflips the residual $\sim 6^\circ$ live orientation error attenuates the loop circles while the donut sweep and altitude are followed.




Fig. 3. SW live-edge (deployed) position, velocity, and orientation error vs. time for slow racing (top), fast racing (middle), and backflips (bottom). Fast-racing position error accumulates monotonically with trajectory length; velocity and orientation errors spike at high-speed, high-dynamics sections; backflip errors peak *during* the flips, where the radar measurement model degrades (Section VIII-C).

The default RANSAC prefilter (Section V-B) buys this same-order position, its clearest single demonstration. A single-chip radar's small virtual aperture gives coarse elevation resolution, and a minority of returns carry an elevation-correlated Doppler bias. Doer's reve rejects these with 3D RANSAC (inlier 0.15 m/s); a Huber-robust least squares only down-weights them. Measured against the ground-truth attitude, the per-frame vertical ego-velocity bias is -0.06 to -0.16 m/s under a Huber-only front-end but at most 0.10 m/s in magnitude under RANSAC. With the prefilter *disabled*, our

vertical channel accumulates a large systematic error (whole-trajectory absolute trajectory error (ATE) 9.6 m on flight_1, $12\text{--}15\%$ vertical drift) that the baselines do not exhibit (theirs stays $\sim 0.2\text{--}0.3 \text{ m}$); the RANSAC prefilter removes it at the source on all four flights, cutting whole-trajectory ATE to $0.24\text{--}0.76 \text{ m}$ (Table IV(b)) and bringing position back into the baselines' range. Their vertical accuracy is not altitude aiding either: disabling the barometer in ekf-yrio shifts its vertical drift by $\leq 0.04 \text{ m}$ (stock $\sigma_{\text{altimeter}} = 5 \text{ m}$ makes its contribution near-negligible), so they too estimate vertical from radar and

IMU alone, by the same outlier-rejection mechanism. The single-chip elevation bias is therefore real but not fundamental: an outlier-rejection gap, closed by the same RANSAC stage the baselines use. Under whole-trajectory Umeyama alignment (Table IV(b)) all three systems agree to ≤ 1 m, consistent with the baselines' published full-alignment figures (they are faithfully reproduced) and confirming that our default front-end matches them.

Their filters on our data. The reverse port is a *sensor/stack-coupling characterization*, not a head-to-head comparison; it is regime-dependent, and we report the measured outcome rather than a verdict. We deliberately keep their *stock* innovation gate and process noise rather than re-tuning them to our platform: the question here is out-of-the-box coupling, and any re-tune would confound it with our own parameter search. Here (the reverse direction) we port their ekf-rio and its multi-radar generalization x-rio [23] (run single-radar) onto our data; ekf-yrio was the comparator in the forward direction (our solver on their data) above. On the slow-racing flight both ekf-rio and x-rio fail: their $\chi^2(3)$ innovation gate rejects 99.7% of the radar ego-velocity updates (301 of 302), so the filter reverts to IMU dead-reckoning and runs essentially unconstrained (> 2 km of error for both, almost entirely vertical, $\approx \frac{1}{2}gt^2$). On the faster flight ekf-rio instead survives, degrading to 1.4 m (3.1% drift, 9.0°): its larger translational excitation produces radar ego-velocities that clear the gate (only 64% rejected), enough to constrain the filter. The radar information admitted thus scales with how strongly motion excites the sparse Doppler; the common cause is a stack-level mismatch, not a defect of their method: our best-velocity firmware yields only ~ 13 points per frame (their RANSAC/ σ -gates expect order-100 clouds, and ego-velocity estimation fails on roughly half of all scans), there is no hardware radar trigger, and our racing-frame vibration exceeds their process-noise range (measured in-flight gyro PSD ≈ 2.1 vs. their $\approx 0.2^\circ/\text{s}/\sqrt{\text{Hz}}$ default), making propagation overconfident so the gate rejects the survivors. This is a design-space split (Section IX), not a defect of either method: their stack runs near sensor rate on dense clouds, ours spends 0.35–0.70 s per window on sparse Doppler, an accuracy-versus-latency trade. This reverse direction thus measures how their stock filters *couple* to an out-of-distribution sensor stack, not the relative quality of the two estimators: any filter starves on ~ 13 -point clouds it was not designed for, so we report it as a sensitivity characterization, not a head-to-head result.

F. Output Covariance Consistency (NEES)

A factor source for a larger fusion graph must report *honest* covariance. Each window we extract the live-edge joint covariance of $(\mathbf{v}_w, \mathbf{R})$ from the converged problem (`ceres::Covariance`, active support control points, tangent space, marginalization prior included) and compute the NEES against the *same* ground truth the headline uses (Viconfed velocity, slerp attitude), after masking windows at MoCap dropouts/glitches. For a consistent estimator each 3-DoF block is χ^2_3 with *mean* 3; we report the per-window mean and, as its spread depends on the window count N , the 95% interval for

TABLE IV
CROSS-VALIDATION ON THE PUBLIC ICINS-2021 FLIGHT DATASETS [20], ALL SYSTEMS UNDER OUR CAUSAL PROTOCOL

(a) drift [% path] / velocity [m/s] / orientation [°]								
Flight (path)	ekf-rio			ekf-yrio			Ours	
1 (142 m)	1.0 / 0.20 / 10.1	0.3 / 0.08 / 0.70			0.6 / 0.08 / 0.51			
2 (45 m)	0.6 / 0.08 / 1.59	0.4 / 0.08 / 0.91			1.1 / 0.08 / 0.94			
3 (147 m)	0.3 / 0.08 / 1.07	0.3 / 0.08 / 1.09			1.3 / 0.07 / 0.52			
4 (78 m)	0.5 / 0.09 / 3.67	0.3 / 0.07 / 1.41			1.2 / 0.07 / 1.13			
(b) Decomposition (systems: ekf-rio / ekf-yrio / Ours)								
Fl.	Roll/pitch [°]			Our drift [%]		Whole ATE [m]		
	rio	yrio	Ours	hor.	ver.	rio	yrio	Ours
1	0.24	0.25	0.25	0.16	0.63	0.87	0.23	0.46
2	0.57	0.59	0.55	0.14	1.13	0.13	0.10	0.24
3	0.25	0.26	0.27	0.30	1.28	0.36	0.26	0.76
4	0.74	0.74	0.74	0.26	1.15	0.30	0.15	0.46

Live estimate, start-anchored alignment. (a) headline: heading aiding differs by design (see text). (b) decomposition: *roll/pitch* is the heading-independent attitude RMSE (yaw excluded), so the orientation gap in (a) is yaw, not attitude; *our drift* splits position into horizontal/vertical (both sub-1.3% under the default RANSAC prefilter, Section V-B); *whole-traj.* ATE re-evaluates every system under EVO/KITTI Umeyama alignment (no scale) for cross-checking published figures. Prefilter *disabled* (Huber-only): vertical drift 12–15%, whole-traj. ATE 9.57/2.88/10.86/5.48 m (flights 1–4).

the N -window mean (not the wide per-sample χ^2_3 interval). On the racing flights the covariance is *calibrated* on fast-racing orientation (NEES mean 3.2, $N=20$, inside the N -window 95% interval [2.0, 4.2]) and *moderately conservative* on the other three channels: slow orientation 1.6 ($N=49$), slow velocity 2.0, fast velocity 1.0, all below their lower bounds (slow [2.4, 3.7]); i.e. the weights over-report uncertainty, most strongly for fast-racing velocity ($\sim 3\times$ in variance). Conservative is the safe direction for a fusion factor. These are *calibration checks*, not powered hypothesis tests: consecutive windows are temporally correlated, so the expectation holds only heuristically. On backflips only 6 of 50 windows survive the ground-truth mask (Section VI), too few for a consistency statement; there the covariance is overconfident (orientation NEES mean 38, $\sim 3.5\times$ in standard deviation; velocity 9.6), reported as a caveat. Deployment recipe: inflate the emitted covariance per regime ($1\times$ gentle, $3\times$ aggressive), or calibrate online from innovations.

G. Ablation Studies

Table V summarizes key design choices.

Key observations from Table V. Full-rate IMU (1 kHz) improves position by $2.9\times$ over 200 Hz, motivating the tight bias priors in the solver configuration. Lever-arm correction helps racing (8–11% position) but has negligible effect on backflips: at $\Delta t_o = 0.8$ ms the dense spline absorbs the $|\boldsymbol{\omega} \times \mathbf{t}_{bs}| \approx 1$ m/s term through orientation DoF. Schur marginalization reduces fast-racing settled position by 18% (0.466 m $\rightarrow$ 0.382 m); for slow racing the optimal prior scale is 10^{-7} (near-zero), so both variants give equivalent results; the benefit is captured in Table VI instead. Window duration has little effect on either racing regime: slow stays at live 0.30–0.34 m across 1.5–5 s, and fast gives essentially equal orientation at 2 s and 3 s (settled 2.33° vs 2.34°) with 3 s modestly better on position (live drift 1.15 vs 1.04%). We keep 3 s as the default for that small position gain, trading latency.

TABLE V
SELECTED ABLATION RESULTS (BATCH SOLVER, SLOW_RACING UNLESS NOTED; SW ROWS USE THE DEFAULT 0.3 s STRIDE; BOLD: DEFAULT CONFIGURATION)

Configuration	Pos (m)	Ori (°)
<i>IMU rate (slow_racing batch, no lever arm[†])</i>		
200 Hz	0.525	1.64
1000 Hz (default)	0.183	1.02
<i>Lever-arm correction $\omega \times t_{bs}$ (batch)</i>		
No correction — slow racing	0.183	1.02
— fast racing	0.862	2.73
— backflips	1.814	8.64
With correction — slow	0.169	1.02
— fast	0.768	2.64
— backflips	1.814	8.44
<i>Orientation regularization (backflips batch)</i>		
Min- ω ($\lambda_\omega=0.001$)	1.816	8.62
Min-α, $\lambda_\alpha=0.1$ (default)	1.814	8.44
Min- α , $\lambda_\alpha=1.0$	2.069	83.8
<i>SW window dur. (slow_racing SW, settled / live edge)</i>		
1.5 s	0.292 / 0.299	1.51 / 2.01
3.0 s (default)	0.293 / 0.310	1.53 / 1.97
5.0 s	0.313 / 0.341	1.49 / 2.03
<i>SW window dur. (fast_racing SW, settled / live edge)</i>		
2.0 s	0.455 / 0.518	2.33 / 2.95
3.0 s (default)	0.382 / 0.469	2.34 / 2.90
<i>Marginalization (fast_racing SW, settled)[‡]</i>		
None (boundary prior only)	0.466	2.68
Schur complement	0.382	2.34

Batch rows (IMU-rate, lever-arm, ori-reg) isolate one design choice at a fixed pre-RANSAC operating point; the RANSAC default shifts batch absolutes (slow batch 0.169→0.198 m / 1.02→1.1°, fast 0.768→0.74 m, the Section VII batch figures) without changing orderings. The SW rows (window duration, marginalization) are under the RANSAC default. All runs use the sensor-only init (Section V-D); the bias seed shifts batch absolutes comparably, again without changing orderings. [†]Lever arm disabled for isolation; see Lever-arm rows for the current default (0.169 m/1.02°). [‡]Marginalization shown for fast_racing; on slow_racing the near-zero optimal prior makes Schur and boundary-only nearly equivalent.

VIII. LIMITATIONS AND NEGATIVE RESULTS

A. Statistical Scope

Each flight regime is evaluated on a *single* representative recording, so the headline numbers carry no run-to-run error bars. The held-out flights (Section VII-D) and the cross-platform validation (Section VII-E) broaden the evidence, but they do not substitute for repeated trials. Reporting mean $\pm$ std over several recordings per regime is the natural strengthening, and is left to future data collection.

B. Heading Sensor Requirement

Yaw is a gauge freedom for Doppler: rotating the entire trajectory about the world gravity axis leaves all radial velocity measurements unchanged, so absolute heading cannot be recovered from radar and IMU alone. A magnetometer resolves this and is standard on flight platforms; here MoCap yaw serves that role. Because that MoCap yaw is noise- and bias-free (unlike a real magnetometer), the heading-aided orientation numbers are an *upper bound*; we bound the real-sensor gap two ways. The dominant magnetometer error is a hard-iron offset: a constant vector in the *body* frame (often 5–15° uncalibrated). In near-level flight its tilt-compensated

horizontal projection is essentially time-constant, hence a heading *gauge* that our start-anchored alignment absorbs, leaving drift, velocity, and orientation RMSE unchanged to first order however large the offset. Only the *time-varying* residual costs: perturbing the prior with a 2° slowly-varying drift plus 1.5° noise (three seeds) degrades live orientation by up to $\sim 1.5^\circ$ on slow racing (2.1–3.5° from 1.97°; the estimate tracks the prior) and $\sim 0.4^\circ$ on fast racing, and a λ_ψ sweep shows this is the sensor's floor, not over-trust ($\lambda_\psi=10$ near-optimal). A throttle-correlated real magnetometer, validated against ground truth, is future work. The heading-*independent* roll/pitch decomposition (Table IV(b)) and the unaided-yaw check (Section VII-E) isolate attitude from this idealization.

C. Extreme Dynamics: Backflips

Batch. The batch solver handles backflips only with $\Delta t_o=0.8$ ms (the *batch* density, $10\times$ denser than the 8 ms sliding-window grid of Table II; 22,630 knots over 18 s) and locked extrinsics (free pitch drifts to $+18^\circ$ due to dense under-constrained orientation DoF). At this density the batch orientation is *bistable* (detailed below): under small timestamp perturbations it oscillates between $\sim 8^\circ$ and $\sim 25^\circ$ regimes, so we do not report a single batch figure and take the sliding-window result as the trustworthy one.

Sliding window. The fixed-lag smoother initially performed poorly on backflips: with the inconsistent prior and a (silently inactive) position anchor, the best result was 2.53 m/10.8° settled. Three fixes account for the working result. (i) The consistent unscaled prior (Section V-E) restores inter-window information transfer. (ii) A soft position anchor to the dead-reckoned initialization ($\lambda_{\text{pos-init}}=0.5$ per CP) bounds the position random walk that sparse, rate-inflated radar cannot prevent within a 3 s window (without it position diverges to 8 m; on racing bags the same anchor is *harmful*: a flip-regime rescue, not a general ingredient). (iii) The dynamics-adaptive weighting law (Section V-B), in particular rate-adaptive gyroscope trust, which on this bag alone is worth $\sim 3^\circ$ of orientation at zero position cost, plus the flip-regime Doppler correction $b=-1.5$ m/s. Together these reach **1.67 m/5.0° settled and 1.55 m/6.4° live at ~ 0.6 s per window** (Table III; Fig. 2). The default RANSAC prefilter (Section V-B) is neutral here (backflip frames are already sparse and mostly fall below its five-return floor, so they bypass it), and the small change from the Huber front-end (1.51 $\rightarrow$ 1.55 m live) is within run-to-run noise. The trust allocation is fully automated (the dynamics-adaptive law); the two rescue terms stay manual because b awaits a better measurement model (per-chirp timestamps) and the tether awaits enough flights to validate an observability-driven adaptive rule.

Why $\Delta t_o=8$ ms in the window. At the batch 0.8 ms a 3 s window is rank-deficient ($\sim 11,250$ orientation DoF vs $\sim 9,000$ gyroscope constraints) and the solver freezes near its initialization; 8 ms is well-conditioned (62.5 Hz bandwidth covers the 10–20 Hz flip dynamics). This same over-dense parameterization is the source of the batch bistability noted above: the final cost *anti-correlates* with accuracy (the underconstrained problem overfits the wrong global orientation), so single-run batch comparisons at this density are not trustworthy.

Attributing the z-axis failure, and what b really is. A visual trajectory check (numeric RMSE conceals it) showed the SW estimate sinking in altitude. Disabling the position anchor exposes the raw behaviour: a near-constant -0.57 m/s world- z drift (-8 m over the flight). Sweeping the per-point correction $v_{\text{corr}} = v - b u_{s,z}$ produces a linear family of z -error curves and large negative b flattens the drift, so we initially interpreted b as the antenna-fixed elevation bias of the array. A joint calibration experiment *disproves* this: (i) with pitch *locked* at its measured value, level-flight racing degrades catastrophically under any explicit b (fast-racing live position $0.39 \rightarrow 3.2 \rightarrow 6.6$ m for $b=0, -0.5, -1.0$), whereas a true antenna-fixed bias would be tolerated; and (ii) on backflips the optimum grows monotonically past the WLS-measured -0.5 m/s (through -1.5 to -2.0) into physically implausible territory. We conclude b is a *flight-regime error proxy*: it absorbs an attitude-correlated Doppler systematic during flips (most likely intra-frame motion distortion) that happens to project like an elevation term, and we cap it at $b=-1.5$ rather than chase the monotone trend without a mechanism. Relatedly, extrinsic pitch is a *plateau*: over a 3×3 grid over (locked pitch, b) the RMSE is indistinguishable across 26.5 – 28.5° on backflips, and locking at 27.5° versus online self-calibration is RMSE-neutral on racing. We freeze it rather than estimate it because free pitch is only weakly observable and *init-dependent* (the converged value tracks the seed: $25^\circ \rightarrow 27^\circ$, $30^\circ \rightarrow 36^\circ$ on both racing bags), and fix it at the physically measured as-built angle (27.5° , inclinometer; Section VII).

Remaining gap. The residual $\sim 6^\circ$ live orientation error during flips attenuates the reconstructed loop circles relative to ground truth. This is not a spline-bandwidth limit (Section VIII-E: denser knots do not help); the residual lever is the radar measurement model, whose intra-frame Doppler smearing during flips is best addressed by *per-chirp timestamps* rather than knot density, which we leave to future work.

D. Prior Scale Sensitivity as a Symptom of Inconsistent Marginalization

Our initial marginalization implementation selected all residuals touching marginalized or boundary variables, which, through the shared bias block, baked every IMU factor of the window into the Schur complement, conditioned on interior states, and double-counted factors re-added in the next window (Section V-E). The resulting overconfident prior had to be deflated by a hand-tuned scale, and the required scale varied by orders of magnitude across flight regimes with a harmful intermediate region. We document this behaviour because the symptom pattern (no universal prior weight, non-monotone sensitivity, robust losses on the prior acting backwards) is a useful *diagnostic for marginalization inconsistency* in sliding-window estimators generally. Table VI shows the sweep measured with the inconsistent prior.

Slow racing is best served by a near-zero prior (10^{-7}): with gentle dynamics, each 3 s window is self-sufficient and additional inter-window coupling is harmful. Fast racing instead needs the prior (2×10^{-4}) for inter-window position continuity; without it (scale = 10^{-7}) fast-racing live position degrades

TABLE VI
PRIOR SCALE SENSITIVITY (SW, LIVE-EDGE RMSE)

scale	Slow racing		Fast racing	
	Pos (m)	Ori ($^\circ$)	Pos (m)	Ori ($^\circ$)
2×10^{-4}	0.543	2.16	0.825	3.64
10^{-5}	—	—	0.942	3.58
10^{-6}	0.442	2.21*	0.951	3.57
10^{-7} (slow default)	0.338	2.08	1.279	3.57
10^{-8}	0.336	2.08	1.344	3.57

*Orientation local maximum: $2.21^\circ > 2.16^\circ$ (2e-4) and 2.08° (1e-7). —: scale not evaluated for this bag.

Measured under the inconsistent-prior diagnostic configuration (predates the RANSAC default); the sensitivity *pattern*, not the absolutes, is the result.

from 0.825 m to 1.279 m. The two bags thus want opposite scales, and the non-monotone slow-racing orientation maximum at 10^{-6} (Table VI) shows no single scale is universally optimal.

The information asymmetry in the inconsistent S (orientation eigenvalues $\sim 5.5 \times 10^{14}$ vs. position $\sim 10^4$) blocks any single scalar, and eigenvalue clipping and Cauchy robustification also fail (Cauchy down-weights the bag that needs the prior most).

Resolution. With the Markov-blanket factor-selection rule of Section V-E, the prior is a true marginal: a single *unscaled* prior (scale = 1) is optimal or near-optimal on all three bags simultaneously, the prior's Mahalanobis residual at each solution is $\mathcal{O}(1)$, and the entire per-mission tuning table becomes obsolete. The fix also improves accuracy: fast-racing settled position improves from 0.727 m (tuned inconsistent prior) to 0.701 m (unscaled consistent prior). Both figures isolate the prior-consistency effect at a fixed pre-RANSAC operating point; the RANSAC default brings the same metric to 0.382 m (Table III). Stride-boundary position jumps in the settled trajectory also disappear (settled velocity RMSE $0.89 \rightarrow 0.21$ m/s on slow racing).

E. Adaptive Knot Density

We tested whether denser orientation knots reduce the residual backflip attitude error; they do not. At $\Delta t_o = 8$ ms the orientation spline already represents the MoCap attitude to 1.28° RMSE *during* the flips, tracking the gyro stream to 0.295 rad/s RMS (flat across $\Delta t_o \in \{16, 8, 4\}$ ms, at the rate-independent vibration floor 0.17 rad/s). The gap is therefore *not* a spline-bandwidth limit: the spline can represent the motion but the optimizer does not rotate to it (open-loop gyro dead-reckoning at 7.8° already beats the 9.2° this fusion reaches before the rate-adaptive gyro trust of Section VIII-C). The lever is measurement weighting, not knot density, *disproving* the adaptive-knot-density hypothesis and obviating the non-uniform-basis extension.

IX. CONCLUSION

Can a single-chip mmWave radar plus IMU provide usable 6-DoF state estimation? Yes, as a *velocity and orientation* source: at the live leading edge, $0.48 \text{ m s}^{-1}/2.0^\circ$ on gentle racing, $0.32 \text{ m s}^{-1}/2.9^\circ$ on aggressive racing (0.35–0.70 s per window on a laptop), confirmed on held-out flights and with

approximately calibrated (NEES) racing orientation covariance. Absolute position is not observable from Doppler+IMU and drifts with distance (0.6–0.9% on racing, 2.9% under tumbling); yaw requires a heading sensor. Within those structural limits, *good results took three things*: (1) a *consistent* fixed-lag smoother, with a Markov-blanket factor-selection rule that makes the marginalization a true marginal, eliminating regime-dependent prior tuning whose symptom pattern we document as a reusable diagnostic; (2) a *dynamics-adaptive trust allocation* across the measurement streams (rate-adaptive gyroscope weight, rate-inflated radar/accelerometer noise, per-point SNR weighting): one weighting law from hover to 10 rad s^{-1} (with per-regime spline grids and two tumbling rescue terms), the velocity–orientation trade made explicit as a measured Pareto curve; and (3) *characterized sensor systematics*, including the finding that the apparent elevation bias under tumbling is a flight-regime error proxy, not an antenna constant, and that the extrinsic pitch is weakly observable, best fixed at its measured value rather than estimated online. Notably, the largest gains came from diagnosis-led re-weighting of existing measurements, not added model capacity: the orientation spline already tracks 10 rad s^{-1} flips at the gyroscope noise floor, disproving our own adaptive-knot-density hypothesis (that denser knots would cut the flip error) without implementing the non-uniform-knot extension.

Open limitations: (1) loop-scale position geometry through flips is bounded by during-flip radar data quality; (2) ground truth itself degrades during flips (MoCap occlusion), bounding what can be measured there; and (3) cross-validation against the EKF-RIO family (Section VII-E) reveals a design-space split: with our RANSAC front-end we reach the same position order of magnitude on their public flights (0.5–1.9 m, trailing the yaw-aided baseline), match their roll/pitch, and lead on heading-aided orientation, while their stock filters, starved by our sparse ~ 13 -point racing clouds, reject up to 99.7% of radar updates. Each stack is engineered for its own regime.

Future directions follow from these results: an incremental, iSAM2-style back-end to turn the demonstrated real-time feasibility (Section VII-B) into onboard margins; per-chirp radar timestamps to target the residual during-flip orientation error our negative results place in the measurement model, not the spline (Section VIII-C); repeated recordings per regime for powered statistics (Section VIII-A). An absolute- z structural anchor (e.g. a floor-plane factor) would make the tumbling-only elevation constant b observable; it should then self-calibrate toward zero, confirming b is a proxy for vertical drift rather than an antenna term, and is left to follow-on work. Since the floor re-enters view intermittently even through the flips, an uncertainty-aware ground model like Go-RIO's [12], gated on the frames that observe it, is a concrete route to this anchor.

REFERENCES

- [1] C. Doer and G. F. Trommer, "An EKF based approach to radar inertial odometry," in *IEEE International Conference on Multisensor Fusion and Integration (MFI)*, 2020, pp. 152–159.
- [2] A. Kramer, C. Stahoviak, A. Santamaria-Navarro, A.-a. Aghamohammadi, and C. Heckman, "Radar-inertial ego-velocity estimation for visually degraded environments," in *IEEE International Conference on Robotics and Automation (ICRA)*, 2020, pp. 5739–5746.
- [3] K. Harlow, H. Jang, T. D. Barfoot, A. Kim, and C. Heckman, "A new wave in robotics: Survey on recent mmWave radar applications in robotics," *IEEE Transactions on Robotics*, vol. 40, pp. 4544–4560, 2024.
- [4] D. Kellner, M. Barjenbruch, J. Klappstein, J. Dickmann, and K. Dietmayer, "Instantaneous ego-motion estimation using doppler radar," in *16th Int. IEEE Conf. Intelligent Transportation Systems (ITSC)*, 2013, pp. 869–874.
- [5] P. Furgale, J. Rehder, and R. Siegwart, "Unified temporal and spatial calibration for multi-sensor systems," in *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2013, pp. 1280–1286.
- [6] V. Usenko, N. Demmel, D. Schubert, J. Stückler, and D. Cremers, "Visual-inertial mapping with non-linear factor recovery," *IEEE Robotics and Automation Letters (RA-L)*, vol. 5, no. 2, pp. 422–429, 2020.
- [7] S. Leutenegger, S. Lynen, M. Bosse, R. Siegwart, and P. Furgale, "Keyframe-based visual-inertial odometry using nonlinear optimization," *International Journal of Robotics Research (IJRR)*, vol. 34, no. 3, pp. 314–334, 2015.
- [8] T. Qin, P. Li, and S. Shen, "VINS-Mono: A robust and versatile monocular visual-inertial state estimator," *IEEE Transactions on Robotics*, vol. 34, no. 4, pp. 1004–1020, 2018.
- [9] Y. Z. Ng, B. Choi, R. Tan, and L. Heng, "Continuous-time radar-inertial odometry for automotive radars," in *IEEE/RSJ Int. Conf. Intelligent Robots and Systems (IROS)*, 2021, pp. 323–330, arXiv:2201.02437.
- [10] K. Burnett, A. P. Schoellig, and T. D. Barfoot, "Continuous-time radar-inertial and lidar-inertial odometry using a gaussian process motion prior," *IEEE Transactions on Robotics*, vol. 41, pp. 1059–1076, 2025, arXiv:2402.06174.
- [11] V.-J. Štirinja, L. Petrović, J. Peršić, I. Marković, and I. Petrović, "Radar-inertial odometry with online spatio-temporal calibration via continuous-time IMU modeling," 2026, arXiv:2603.19958.
- [12] W. Yang, H. Jang, and A. Kim, "Ground-optimized 4d radar-inertial odometry via continuous velocity integration using gaussian process," 2025, arXiv:2502.08093.
- [13] Q. Huang, Y. Liang, Z. Qiao, S. Shen, and H. Yin, "Less is more: Physical-enhanced radar-inertial odometry," in *IEEE Int. Conf. Robotics and Automation (ICRA)*, 2024, pp. 15 966–15 972, arXiv:2402.02200.
- [14] J. Michalczyk, R. Jung, and S. Weiss, "Tightly-coupled EKF-based radar-inertial odometry," in *IEEE/RSJ Int. Conf. Intelligent Robots and Systems (IROS)*, 2022, pp. 12 336–12 343.
- [15] J. Michalczyk, J. Quell, F. Steidle, M. G. Müller, and S. Weiss, "Tightly-coupled factor graph formulation for radar-inertial odometry," in *IEEE/RSJ Int. Conf. Intelligent Robots and Systems (IROS)*, 2024, pp. 3364–3370.
- [16] R. Girod, M. Hauswirth, P. Pfreundschuh, M. Biasio, and R. Siegwart, "A robust baro-radar-inertial odometry M-estimator for multicopter navigation in cities and forests," in *IEEE Int. Conf. Multisensor Fusion and Integration (MFI)*, 2024, pp. 1–8, arXiv:2408.05764.
- [17] T.-C. Dong-Si and A. I. Mourikis, "Consistency analysis for sliding-window visual odometry," in *IEEE Int. Conf. Robotics and Automation (ICRA)*, 2012, pp. 5202–5209.
- [18] J. Lv, X. Lang, J. Xu, M. Wang, Y. Liu, and X. Zuo, "Continuous-time fixed-lag smoothing for lidar-inertial-camera SLAM," *IEEE/ASME Transactions on Mechatronics*, vol. 28, no. 4, pp. 2259–2270, 2023, arXiv:2302.07456.
- [19] J. Quenzel and S. Behnke, "LIO-MARS: Non-uniform continuous-time trajectories for real-time lidar-inertial-odometry," 2025, arXiv:2511.13985.
- [20] C. Doer and G. F. Trommer, "Yaw aided radar inertial odometry using manhattan world assumptions," in *28th Saint Petersburg Int. Conf. Integrated Navigation Systems (ICINS)*, 2021, pp. 1–9.
- [21] C. Sommer, V. Usenko, D. Schubert, N. Demmel, and D. Cremers, "Efficient derivative computation for cumulative B-splines on Lie groups," in *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020, pp. 11 148–11 156.
- [22] H. Martiros, A. Miller, N. Bucki, B. Solliday, R. Kennedy, J. Zhu, T. Dang, D. Pattison, H. Zheng, T. Tomic, P. Henry, G. Cross, J. Vander-Mey, A. Sun, S. Wang, and K. Holtz, "SymForce: Symbolic computation and code generation for robotics," in *Robotics: Science and Systems (RSS)*, 2022.
- [23] C. Doer and G. F. Trommer, "x-RIO: Radar inertial odometry with multiple radar sensors and yaw aiding," *Gyroscopy and Navigation*, vol. 12, no. 4, pp. 329–339, 2021.